

MINI PROJECT

PROJECT TOPIC : Prediction and Analysis on Diabetes Dataset

Project Description :

The system will allow us to predict if the patient has diabetes on the basis of certain diagnostic measures available in the dataset. The different steps involved in EDA include: 1.Data Collection, 2.Data Cleaning and 3.Data Visualization. This project first conducts Exploratory Data Analysis (EDA) and data visualization on the diabetes dataset and then predict the diabetes.

Exploratory Data Analysis (EDA)

1. Descriptive statistics

Attribute type, Class distribution, Mean, Standard Deviation, Median, Quartile, Skewness, Correlation

2. Data Visualization

Histogram plot

Density plot

Box and Whisker plot

Bar plot

Missing data map

Pair-wise correlation plot

Prediction on Diabetes

We compare the performance for the following classifiers:

1. Logistic Regression

Logistic regression is basically a supervised classification algorithm. In a classification problem, the target variable(or output), y , can take only discrete values for given set of features(or inputs), X .

2. Support Vector Machine (SVM)

“Support Vector Machine” (SVM) is a supervised machine learning algorithm which can be used for both classification or regression challenges. However, it is mostly used in classification problems. In the SVM algorithm, we plot each data item as a point in n -dimensional space (where n is number of features you have) with the value of each feature being the value of a particular coordinate. Then, we perform classification by finding the hyper-plane that differentiates the two classes very well

3. random Forest

The random forest algorithm works by aggregating the predictions made by multiple decision trees of varying depth. Every decision tree in the forest is trained on a subset of the dataset called the bootstrapped dataset.

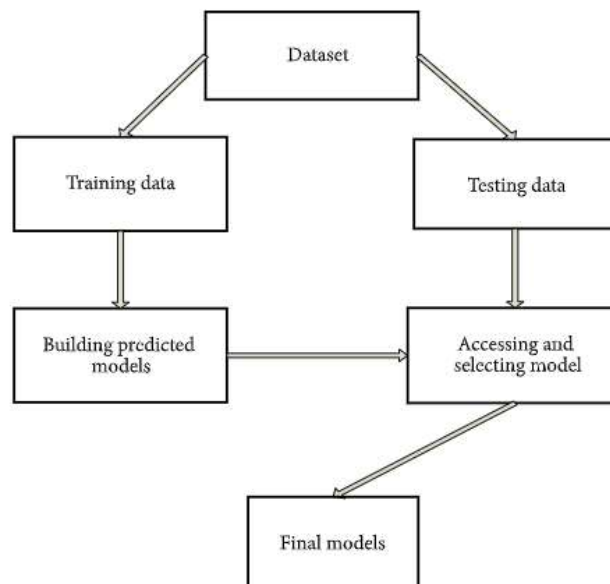
Dataset:

The dataset used in this study is downloaded from:

<http://biostat.mc.vanderbilt.edu/wiki/Main/DataSets?CGISESSID=10713f6d891653ddcbb7ddb9cffb79>

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S
1	id	chol	stab.glu	hdl	ratio	glyhb	location	age	gender	height	weight	frame	bp.1s	bp.1d	bp.2s	bp.2d	waist	hip	time.ppn
2	1000	203	82	56	3.6	4.31	Buckingha	46	female	62	121	medium	118	59			29	38	720
3	1001	165	97	24	6.9	4.44	Buckingha	29	female	64	218	large	112	68			46	48	360
4	1002	228	92	37	6.2	4.64	Buckingha	58	female	61	256	large	190	92	185	92	49	57	180
5	1003	78	93	12	6.5	4.63	Buckingha	67	male	67	119	large	110	50			33	38	480
6	1005	249	90	28	8.9	7.72	Buckingha	64	male	68	183	medium	138	80			44	41	300
7	1008	248	94	69	3.6	4.81	Buckingha	34	male	71	190	large	132	86			36	42	155
8	1011	195	92	41	4.8	4.84	Buckingha	30	male	69	191	medium	161	112	161	112	46	49	720
9	1015	227	75	44	5.2	3.94	Buckingha	37	male	59	170	medium					34	39	1020
10	1016	177	87	49	3.6	4.84	Buckingha	45	male	69	166	large	160	80	128	86	34	40	300
11	1022	263	89	40	6.6	5.78	Buckingha	55	female	63	262	small	108	72			45	50	240
12	1024	242	82	54	4.5	4.77	Louisa	60	female	65	156	medium	130	90	130	90	39	45	300
13	1029	215	128	34	6.3	4.97	Louisa	38	female	58	195	medium	102	68			42	50	90
14	1030	238	75	36	6.6	4.47	Louisa	27	female	60	170	medium	130	80			35	41	720
15	1031	183	79	46	4	4.59	Louisa	40	female	59	165	medium					37	43	60
16	1035	191	76	30	6.4	4.67	Louisa	36	male	69	183	medium	100	66			36	40	225
17	1036	213	83	47	4.5	3.41	Louisa	33	female	65	157	medium	130	90	120	96	37	41	240
18	1037	255	78	38	6.7	4.33	Louisa	50	female	65	183	medium	130	100			37	43	180
19	1041	230	112	64	3.6	4.53	Louisa	20	male	67	159	medium	100	90			31	39	1440
20	1045	194	81	36	5.4	5.28	Louisa	36	male	64	126	medium	110	76			30	34	120
21	1250	196	206	41	4.8	11.24	Buckingha	62	female	65	196	large	178	90			46	51	540
22	1252	186	97	50	3.7	6.49	Buckingha	70	male	67	178	large	148	88	148	84	42	41	1020

Our research dataset is divided into two parts; two-thirds of the data is used as a training set, and one-third of the dataset is defined as a testing set to evaluate the performance of several classifiers. All classifiers were fitted to the same training and testing data. The specific process is: .



Libraries:

- **randomforest**

randomForest implements Breiman's random forest algorithm (based on Breiman and Cutler's original Fortran code) for classification and regression. It can also be used in unsupervised mode for assessing proximities among data points.

- **Caret**

The caret package (short for Classification And REgression Training) is a set of functions that attempt to streamline the process for creating predictive models. The package contains tools for:

- data splitting
- pre-processing
- feature selection
- model tuning using resampling
- variable importance estimation
- as well as other functionality.
-

For EDA and Visualization Steps performed:

1. Load libraries
2. Load the data
3. Clean the data
4. Delete columns 15 and 16 due to many missing values and column 1 (id), column 7 (location) because they contain no useful information
5. Remove the row with missing values
6. Encode the class label (column 5): Glycosolated hemoglobin > 7.0 is taken as a positive diagnosis of diabetes
7. Encode the categorical data (column-7 gender)
8. Encode the categorical data (column-10 frame)
9. Descriptive statistics
10. Display the first 20 rows
11. Display the dimensions of the dataset
12. List types for each attribute
13. Distribution of the class labels
14. Summarize the dataset
15. Standard Deviations for the non-categorical columns
16. Perrom other operations like: Skewness, Correlation, and Visualization Plots

Predcition and Classification Steps performed:

1. Load libraries
2. Load the data
3. Clean the data
4. Delete columns 15 and 16 due to many missing values and column 1 (id), column 7 (location) because they contain no useful information
5. Remove the row with missing values
6. Encode the class label (column 5): Glycosolated hemoglobin > 7.0 is taken as a positive diagnosis of diabetes
7. Encode the categorical data (column-7 gender)

8. Encode the categorical data (column-10 frame)
9. Split the data into training and validation sets
10. Comparison among different classifiers
 1. Logistic Regression
 2. Support Vector Machine
 3. random forest
 4. Parameter tuning via grid search for random forest
11. Make predictions on the validation sets
12. Confusion matrix
13. Save the final classifier model into disk
14. Load the model from the disk
15. Make predictions using the loaded model
15. Make predictions using the loaded model

Output:

```
# Load libraries
```

```
library(randomForest)
```

```
library(caret)
```

```
# load the data
```

```
> filename="C:\\Users\\Tejal Chaudhary\\Desktop\\Exploratory-Data-Analysis  
-and-Prediction-on-Diabetes-Dataset-using-R-master\\diabetes_data.csv"
```

```
> datasetRaw = read.csv(filename)
```

```
> print(head(datasetRaw))
```

	id	cho1	stab	glu	hdl	ratio	glyhb	location	age	gender	height	weight
	frame	bp.1s	bp.1d	bp.2s	bp.2d							
1	1000	203		82	56	3.6	4.31	Buckingham	46	female	62	121 m
	edium	118	59		NA	NA						
2	1001	165		97	24	6.9	4.44	Buckingham	29	female	64	218
	large	112	68		NA	NA						
3	1002	228		92	37	6.2	4.64	Buckingham	58	female	61	256
	large	190	92	185	92							
4	1003	78		93	12	6.5	4.63	Buckingham	67	male	67	119
	large	110	50		NA	NA						
5	1005	249		90	28	8.9	7.72	Buckingham	64	male	68	183 m
	edium	138	80		NA	NA						
6	1008	248		94	69	3.6	4.81	Buckingham	34	male	71	190
	large	132	86		NA	NA						

```
waist hip time.ppn
```

1	29	38	720
2	46	48	360
3	49	57	180
4	33	38	480
5	44	41	300
6	36	42	195

```
# clean the data
```

```
> numColumns = dim(datasetRaw)[2]
```

```
> vector_NAs = rep(0, numColumns)
```

```
> for (i in 1:numColumns) {  
+   vector_NAs[i] = sum(is.na(datasetRaw[,i]))
```

```
+  
+ }
```

```
> print("The missing values in each column:")
```

```
[1] "The missing values in each column:"
```

```
> print(vector_NAs)
```

```
[1] 0 1 0 1 1 13 0 0 0 5 1 0 5 5 262 262 2  
2 3
```

```
# delete columns 15 and 16 due to many missing values
```

```

# delete column 1 (id), column 7 (location) because they contain no useful information
> print(dim(dataset))
[1] 403 15
# remove the row with missing values
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
[1] 375 15
> print(head(dataset))
  chol stab.glu hdl ratio glyhb age gender height weight frame bp.1s bp.1
d waist hip time.ppn
1 203 29 38 82 56 3.6 4.31 46 female 62 121 medium 118 5
9 165 46 48 97 24 6.9 4.44 29 female 64 218 large 112 6
2 228 49 57 92 37 6.2 4.64 58 female 61 256 large 190 9
2 78 33 38 93 12 6.5 4.63 67 male 67 119 large 110 5
0 249 44 41 90 28 8.9 7.72 64 male 68 183 medium 138 8
6 248 36 42 94 69 3.6 4.81 34 male 71 190 large 132 8
6 195 42 49 92 41 4.8 0 30 1 69 191 1 161 11
2 177 49 87 49 3.6 0 45 1 69 166 2 160 8
0 34 40 300 6.6 0 55 0 63 202 0 108 7
2 45 50 240 4.5 0 60 0 65 156 1 130 9
0 39 45 300 6.3 0 38 0 58 195 1 102 6
8 42 50 90 6.6 0 27 0 60 170 1 130 8
13 238 35 41 75 36 6.4 0 36 1 69 183 1 100 6
6 36 40 225 4.5 0 33 0 65 157 1 130 9
16 213 37 41 83 47 6.7 0 50 0 65 183 1 130 10
0 37 43 180

```

18	230		112	64	3.6	0	20	1	67	159	1	100	9
0	31	39		1440									
19	194		81	36	5.4	0	36	1	64	126	1	110	7
6	30	34		120									
20	196		206	41	4.8	1	62	0	65	196	2	178	9
0	46	51		540									
21	186		97	50	3.7	0	70	1	67	178	2	148	8
8	42	41		1020									
22	234		65	76	3.1	0	47	1	67	230	2	137	10
0	45	46		480									

```
# display the dimensions of the dataset
> print(dim(dataset))
[1] 375 15
> print(head(dataset))
  chol stab.glu hdl ratio glyhb age gender height weight frame bp.1s bp.1
d waist hip time.ppn
1 203 29 38 82 56 3.6 4.31 46 female 62 121 medium 118 5
9 29 38 720
2 165 46 48 97 24 6.9 4.44 29 female 64 218 large 112 6
8 46 48 360
3 228 49 57 92 37 6.2 4.64 58 female 61 256 large 190 9
2 49 57 180
4 78 33 38 93 12 6.5 4.63 67 male 67 119 large 110 5
0 33 38 480
5 249 44 41 90 28 8.9 7.72 64 male 68 183 medium 138 8
0 44 41 300
6 248 36 42 94 69 3.6 4.81 34 male 71 190 large 132 8
6 36 42 195
> # encode the class label (column 5): Glycosolated hemoglobin > 7.0 is t
aken as a positive diagnosis of diabetes.
> dataset[,5] = ifelse(dataset[,5] >= 7.0, 1, 0)
> dataset[,5] = factor(dataset[,5])
> # encode the categorical data (column-7 gender)
> dataset[,7] = ifelse(dataset[,7] == "female", 0, 1)
> dataset[,7] = factor(dataset[,7])
> # encode the categorical data (column-10 frame)
> dataset[,10] = ifelse(dataset[,10] == "small", 0, ifelse(dataset[,10] ==
"medium", 1,2) )
> dataset[,10] = factor(dataset[,10])
> # Descriptive statistics %%%%%%%%%
> # display the first 20 rows
> print(head(dataset, n=20))
  chol stab.glu hdl ratio glyhb age gender height weight frame bp.1s bp.1
d waist hip time.ppn
1 203 29 38 82 56 3.6 0 46 0 62 121 1 118 5
9 29 38 720
2 165 46 48 97 24 6.9 0 29 0 64 218 2 112 6
8 46 48 360
3 228 49 57 92 37 6.2 0 58 0 61 256 2 190 9
2 49 57 180
4 78 33 38 93 12 6.5 0 67 1 67 119 2 110 5
0 33 38 480
5 249 44 41 90 28 8.9 1 64 1 68 183 1 138 8
0 44 41 300
6 248 36 42 94 69 3.6 0 34 1 71 190 2 132 8
6 36 42 195
7 195 46 49 92 41 4.8 0 30 1 69 191 1 161 11
2 46 49 720
9 177 34 40 87 49 3.6 0 45 1 69 166 2 160 8
0 34 40 300
10 263 45 50 89 40 6.6 0 55 0 63 202 0 108 7
2 45 50 240
11 242 39 45 82 54 4.5 0 60 0 65 156 1 130 9
0 39 45 300
12 215 42 50 128 34 6.3 0 38 0 58 195 1 102 6
8 42 50 90
13 238 75 36 6.6 0 27 0 60 170 1 130 8
0 35 41 720
```

```

15 191      76 30    6.4    0 36      1    69    183    1 100    6
6   36 40      225
16 213      83 47    4.5    0 33      0    65    157    1 130    9
0   37 41      240
17 255      78 38    6.7    0 50      0    65    183    1 130   10
0   37 43      180
18 230      112 64   3.6    0 20      1    67    159    1 100    9
0   31 39      1440
19 194      81 36    5.4    0 36      1    64    126    1 110    7
6   30 34      120
20 196      206 41   4.8    1 62      0    65    196    2 178    9
0   46 51      540
21 186      97 50    3.7    0 70      1    67    178    2 148    8
8   42 41      1020
22 234      65 76    3.1    0 47      1    67    230    2 137   10
0   45 46      480

```

```
> # display the dimensions of the dataset
```

```
> print(dim(dataset))
```

```
[1] 375 15
```

```
> # list types for each attribute
```

```
> print(sapply(dataset,class))
```

```

chol stab.glu hdl ratio glyhb age gender h
eight weight frame bp.ls
"integer" "integer" "integer" "numeric" "factor" "integer" "factor" "int
eger" "integer" "factor" "integer"
bp.lid waist hip time.ppn
"integer" "integer" "integer" "integer"

```

```
> # distribution of the class labels
```

```
> y = dataset$glyhb
```

```
> print(cbind(freq=table(y),percentage=prop.table(table(y))*100))
```

```
freq percentage
```

```
0 317 84.53333
```

```
1 58 15.46667
```

```
> print(table(y))
```

```
y
```

```
0 1
```

```
317 58
```

```
> print(table(y)/length(y))
```

```
y
```

```
0 1
```

```
0.8453333 0.1546667
```

```
> # summarize the dataset
```

```
> print(summary(dataset))
```

```

chol stab.glu hdl ratio glyhb
age gender height
Min. : 78.0 Min. : 48.0 Min. : 12.00 Min. : 1.500 0:317
Min. :19.00 0:220 Min. :52
1st Qu.:179.0 1st Qu.: 81.0 1st Qu.: 38.00 1st Qu.: 3.200 1: 58
1st Qu.:34.00 1:155 1st Qu.:63
Median :204.0 Median : 90.0 Median : 46.00 Median : 4.200
Median :45.00 Median :66
Mean :207.6 Mean :107.6 Mean : 50.43 Mean : 4.525
Mean :46.98 Mean :66
3rd Qu.:229.5 3rd Qu.:108.5 3rd Qu.: 59.00 3rd Qu.: 5.400
3rd Qu.:60.00 3rd Qu.:69
Max. :443.0 Max. :385.0 Max. :120.00 Max. :19.300
Max. :92.00 Max. :76

weight frame bp.ls bp.lid waist
hip time.ppn
Min. : 99.0 0: 98 Min. : 90.0 Min. : 48.00 Min. :26.00
Min. :30.00 Min. : 5
1st Qu.:151.0 1:172 1st Qu.:121.5 1st Qu.: 75.00 1st Qu.:33.00
1st Qu.:39.00 1st Qu.: 90
Median :174.0 2:105 Median :136.0 Median : 82.00 Median :37.00
Median :42.00 Median :240
Mean :177.9 Mean :137.5 Mean : 83.38 Mean :37.95
Mean :43.09 Mean :335
3rd Qu.:200.0 3rd Qu.:148.0 3rd Qu.: 91.00 3rd Qu.:41.50
3rd Qu.:46.00 3rd Qu.: 480

```

```

Max.      :325.0      Max.      :250.0      Max.      :124.00      Max.      :56.00
Max.      :64.00      Max.      :1560
# Standard Deviations for the non-categorical columns
> std=apply(dataset[,-c(5,7,10)],sd)
> print('The standard deviations are:')
[1] "The standard deviations are:"
> print(std)
chol      stab.glu      hdl      ratio      age      height      weight      bp.1s
44.700780  54.082496  17.444346  1.755499  16.661203  3.915210  40.56894
0  23.178154
      bp.1d      waist      hip      time.ppn
13.544167  5.777105  5.642679  309.056806
# Skewness
#The further the distribution of the skew value from zero,
# the larger the skew to the left (negative skew value) or right (positive
skew value).
> library(e1071) # the library for skewness
> skew=apply(dataset[,-c(5,7,10)], 2, skewness)
> print(skew)
      chol      stab.glu      hdl      ratio      age      height      weig
ht      bp.1s
0.97739823  2.69790949  1.21275829  2.24132546  0.30061280  0.02678693  0.748807
75 1.05634395
      bp.1d      waist      hip      time.ppn
0.23310577  0.47060516  0.80724954  1.28077802
# Correlations
> correlations=cor(dataset[,-c(5,7,10)])
> print(correlations)
      chol      stab.glu      hdl      ratio      age
height      1.0000000000  0.17605760  0.18918622  0.47685119  0.25807515 -0.
0726825795
stab.glu      0.1760576046  1.00000000  -0.15845068  0.29353513  0.28820616  0.
0864302351
hdl      0.1891862222  -0.15845068  1.00000000  -0.68325757  0.02624555 -0.
0901427069
ratio      0.4768511857  0.29353513  -0.68325757  1.00000000  0.16952389  0.
0782807237
age      0.2580751516  0.28820616  0.02624555  0.16952389  1.00000000 -0.
0953003330
height      -0.0726825795  0.08643024  -0.09014271  0.07828072  -0.09530033  1.
0000000000
weight      0.0545001654  0.17918946  -0.29760400  0.27690932  -0.06523800  0.
2460901156
bp.1s      0.2084006482  0.16181629  0.03436444  0.11090086  0.45236549 -0.
0470999583
bp.1d      0.1715635556  0.02867443  0.07630061  0.03904864  0.06562279  0.
0383903197
waist      0.1192740454  0.22720469  -0.28499211  0.30894469  0.15068806  0.
0450335672
hip      0.0732717715  0.13333204  -0.23280194  0.20246310  -0.00208207 -0.
1176291370
time.ppn      0.0009236077  -0.05010410  0.06822885  -0.05007620  -0.03962976 -0.
0001877958
      weight      bp.1s      bp.1d      waist      hip
time.ppn
chol      0.05450017  0.20840065  0.17156356  0.11927405  0.07327177  0.00
09236077
stab.glu      0.17918946  0.16181629  0.02867443  0.22720469  0.13333204  -0.05
01040985
hdl      -0.29760400  0.03436444  0.07630061  -0.28499211  -0.23280194  0.06
82288499
ratio      0.27690932  0.11090086  0.03904864  0.30894469  0.20246310  -0.05
00762030
age      -0.06523800  0.45236549  0.06562279  0.15068806  -0.00208207  -0.03
96297574
height      0.24609012  -0.04709996  0.03839032  0.04503357  -0.11762914  -0.00
01877958

```



```


weight      1.00000000  0.08797282  0.17190112  0.85176414  0.82931024 -0.05
77287098
bp.1s       0.08797282  1.00000000  0.61104157  0.20342942  0.14425492 -0.08
14349347
bp.1d       0.17190112  0.61104157  1.00000000  0.17056560  0.15538918 -0.06
69496556
waist        0.85176414  0.20342942  0.17056560  1.00000000  0.83224050 -0.06
47235445
hip         0.82931024  0.14425492  0.15538918  0.83224050  1.00000000 -0.09
43397549
time.ppn    -0.05772871 -0.08143493 -0.06694966 -0.06472354 -0.09433975  1.00
00000000

```

```

# Data visualizations
> dataset_numeric = dataset[,-c(5,7,10)]
> #Histograms
> par(mfrow=c(3,4)) # put four figures in a row (2*4)
> for (i in 1:12) {
+   hist(dataset_numeric[,i],main=names(dataset_numeric)[i])
+ }

```

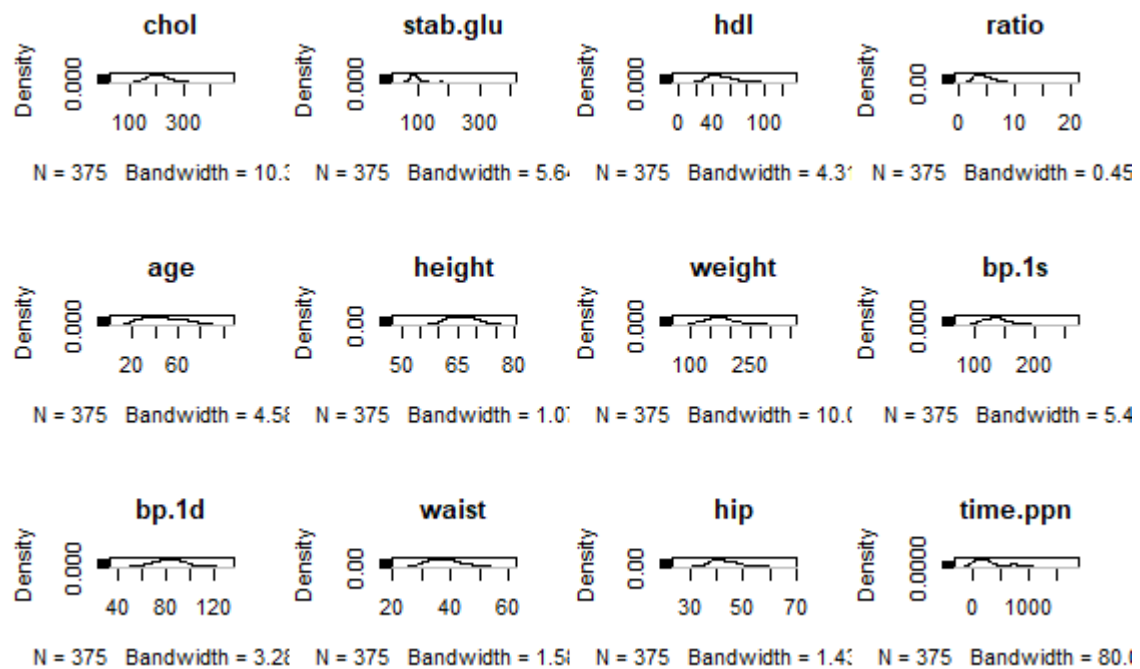


```

#Density Plots
> par(mfrow=c(3,4))
> for(i in 1:12) {
+   plot(density(dataset_numeric[,i]), main=names(dataset_numeric)[i])
+ }


```

```
+ }
```



```
#Box And Whisker Plots
```

```
> par(mfrow=c(3,4))
> for(i in 1:12) {
+   boxplot(dataset_numeric[,i], main=names(dataset_numeric)[i])
+ }
```




```
#Barplots, which is used to count the accuracies for categorical attribute
```

```
s
> dataset_categorical = dataset[,c(5,7,10)]
> par(mfrow=c(1,3))
> for(i in 1:3) {
```

```

+ counts <- table(dataset_categorical[,i]) # get the count for each categorical value
+ name <- names(dataset_categorical)[i]
+ barplot(counts, main=name)
+ }
#Missing data Plot


```



```

> library(Amelia) # library for the function missmap
> par(mfrow=c(1,1))
> missmap(datasetRaw, col=c("red", "grey"), legend=FALSE)

```



```

#Multivariate Visualization
> library(corrplot) # for function corrplot()
corrplot 0.84 loaded
> correlations1=cor(dataset_numeric)

```

```

> print(correlations1)

```

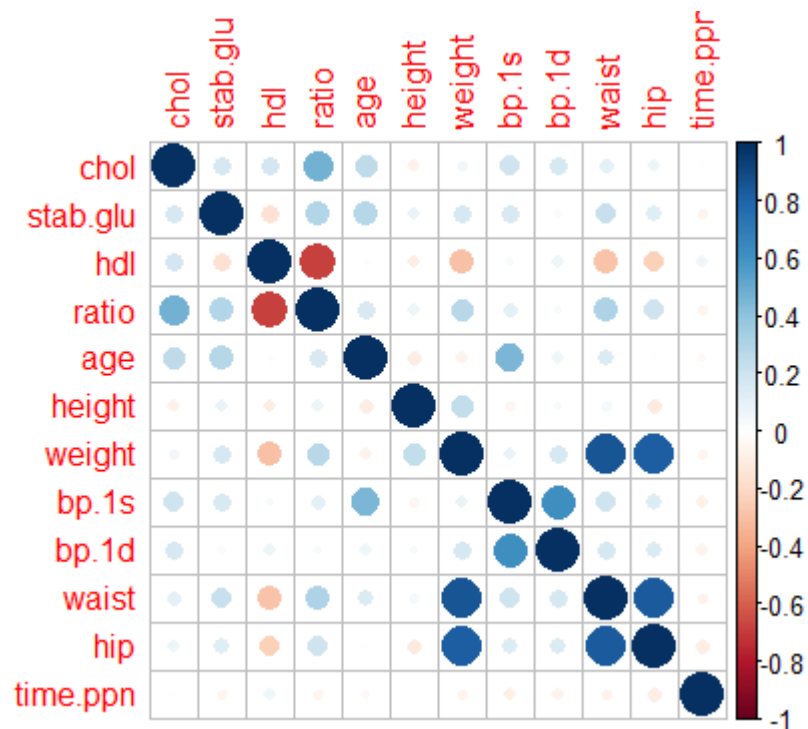
	cho1	stab.glu	hdl	ratio	age
height					
cho1	1.0000000000	0.17605760	0.18918622	0.47685119	0.25807515
0726825795					
stab.glu	0.1760576046	1.00000000	-0.15845068	0.29353513	0.28820616
0864302351					
hdl	0.1891862222	-0.15845068	1.00000000	-0.68325757	0.02624555
0901427069					
ratio	0.4768511857	0.29353513	-0.68325757	1.00000000	0.16952389
0782807237					
age	0.2580751516	0.28820616	0.02624555	0.16952389	1.00000000
0953003330					
height	-0.0726825795	0.08643024	-0.09014271	0.07828072	-0.09530033
0000000000					
weight	0.0545001654	0.17918946	-0.29760400	0.27690932	-0.06523800
2460901156					
bp.1s	0.2084006482	0.16181629	0.03436444	0.11090086	0.45236549
0470999583					
bp.1d	0.1715635556	0.02867443	0.07630061	0.03904864	0.06562279
0383903197					
waist	0.1192740454	0.22720469	-0.28499211	0.30894469	0.15068806
0450335672					
hip	0.0732717715	0.13333204	-0.23280194	0.20246310	-0.00208207
1176291370					
time.ppn	0.0009236077	-0.05010410	0.06822885	-0.05007620	-0.03962976
0001877958					

	weight	bp.1s	bp.1d	waist	hip
time.ppn					
cho1	0.05450017	0.20840065	0.17156356	0.11927405	0.07327177
09236077					
stab.glu	0.17918946	0.16181629	0.02867443	0.22720469	0.13333204
01040985					
hdl	-0.29760400	0.03436444	0.07630061	-0.28499211	-0.23280194
82288499					
ratio	0.27690932	0.11090086	0.03904864	0.30894469	0.20246310
00762030					
age	-0.06523800	0.45236549	0.06562279	0.15068806	-0.00208207
96297574					
height	0.24609012	-0.04709996	0.03839032	0.04503357	-0.11762914
01877958					
weight	1.00000000	0.08797282	0.17190112	0.85176414	0.82931024
77287098					
bp.1s	0.08797282	1.00000000	0.61104157	0.20342942	0.14425492
14349347					
bp.1d	0.17190112	0.61104157	1.00000000	0.17056560	0.15538918
69496556					
waist	0.85176414	0.20342942	0.17056560	1.00000000	0.83224050
47235445					
hip	0.82931024	0.14425492	0.15538918	0.83224050	1.00000000
43397549					
time.ppn	-0.05772871	-0.08143493	-0.06694966	-0.06472354	-0.09433975
00000000					


```

> par(mfrow=c(1,1))
> corrplot(correlations1, methods="circle")

```

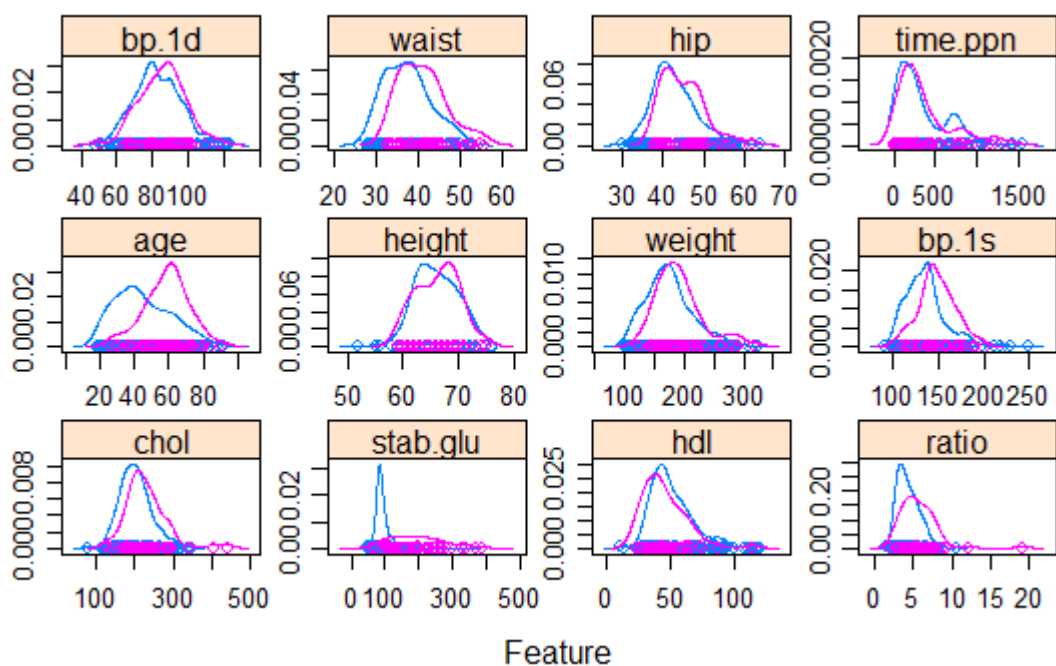


```
# pair-wise scatterplots of the numeric attributes
> par(mfrow=c(1,1))
pairs(dataset_numeric)
> #Scatterplot Matrix By Class (use different color to distinguish different class)
> par(mfrow=c(1,1))
> pairs(dataset_numeric, col=dataset[,5])
```



```
#Density By Class
> library(caret)
> # load the data
> data(iris)
> # density plots for each attribute by class value
> x <- dataset_numeric
> y <- dataset[,5]
> scales <- list(x=list(relation="free"), y=list(relation="free"))
```

```
> par(mfrow=c(1,1))
> featurePlot(x=dataset_numeric, y=dataset[,5], plot="density", scales=scales)
```



```
#Box And Whisker Plots By Class
> featurePlot(x=dataset_numeric, y=dataset[,5], plot="box")
```



```

# Load libraries
> library(randomForest)
> library(caret)

> # load the data
> filename="C:\\Users\\Tejal Chaudhary\\Desktop\\Exploratory-Data-Analysis-and
-Prediction-on-Diabetes-Dataset-using-R-master\\diabetes_data.csv"
> datasetRaw = read.csv(filename)
> print(head(datasetRaw))
  id chol stab.glu hdl ratio glyhb location age gender height weight
frame bp.1s
1 1000 203      82  56   3.6  4.31 Buckingham 46 female     62    121 m
edum 118
2 1001 165      97  24   6.9  4.44 Buckingham 29 female     64    218
large 112
3 1002 228      92  37   6.2  4.64 Buckingham 58 female     61    256
large 190
4 1003 78       93  12   6.5  4.63 Buckingham 67 male       67    119
large 110
5 1005 249      90  28   8.9  7.72 Buckingham 64 male       68    183 m
edum 138
6 1008 248      94  69   3.6  4.81 Buckingham 34 male       71    190
large 132
bp.1d bp.2s bp.2d waist hip time.ppn
1 59 NA NA 29 38 720
2 68 NA NA 46 48 360
3 92 185 92 49 57 180
4 50 NA NA 33 38 480
5 80 NA NA 44 41 300
6 86 NA NA 36 42 195

# clean the data
> numColumns = dim(datasetRaw)[2]
> vector_NAs = rep(0, numColumns)
> for (i in 1:numColumns) {
+   vector_NAs[i] = sum(is.na(datasetRaw[,i]))
+ }
> print("The missing values in each column:")
[1] "The missing values in each column:"
> print(vector_NAs)
[1] 0 1 0 1 1 13 0 0 0 5 1 0 5 5 262 262 2
2 3

# delete columns 15 and 16 due to many missing values
> # delete column 1 (id), column 7 (location) because they contain no usef
ul information
> dataset = datasetRaw[, -c(1,7,15,16)]
> print(dim(dataset))
[1] 403 15

# remove the row with missing values
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
[1] 375 15
> print(head(dataset))
chol stab.glu hdl ratio glyhb age gender height weight frame bp.1s bp.1
d waist hip
1 203 29 38 82 56 3.6 4.31 46 female 62 121 medium 118 5
2 165 46 48 97 24 6.9 4.44 29 female 64 218 large 112 6
3 228 49 57 92 37 6.2 4.64 58 female 61 256 large 190 9
4 78 33 38 93 12 6.5 4.63 67 male 67 119 large 110 5
0

```

```

5 249      90 28 8.9 7.72 64 male 68 183 medium 138 8
0 44 41
6 248      94 69 3.6 4.81 34 male 71 190 large 132 8
6 36 42
time.ppn
1 720
2 360
3 180
4 480
5 300
6 195

```

encode the class label (column 5): Glycosolated hemoglobin > 7.0 is taken as a positive diagnosis of diabetes.

```

> dataset[,5] = ifelse(dataset[,5] >= 7.0, 1, 0)
> dataset[,5] = factor(dataset[,5])
> # encode the categorical data (column-7 gender)
> dataset[,7] = ifelse(dataset[,7] == "female", 0, 1)
> dataset[,7] = factor(dataset[,7])
> # encode the categorical data (column-10 frame)
> dataset[,10] = ifelse(dataset[,10] == "small", 0, ifelse(dataset[,10] == "medium", 1, 2))
> dataset[,10] = factor(dataset[,10])
> # split the data into training and validation sets
> set.seed(7)
> validation_index = createDataPartition(dataset$glyhb, p=0.90, list=FALSE)
> validationData = dataset[-validation_index,]
> trainingData = dataset[validation_index,]
> # comparison among different classifiers
> #1. Logistic Regression
> set.seed(7)
> control.glm = trainControl(method = "cv", number = 5)
> fit.glm = train(glyhb~., data = trainingData, method = "glm", preProc = c("center", "scale"), trControl = control.glm)
> print(fit.glm$results)
  parameter Accuracy      Kappa AccuracySD      KappaSD
1      none 0.9113696 0.6246138 0.03189889 0.1209417

```

#2. Support Vector Machine

```

> set.seed(7)
> control.svmRadial = trainControl(method="cv", number=5)
> fit.svmRadial <- train(glyhb~., data=trainingData, method="svmRadial", metric="Accuracy", preProc=c("center", "scale"), trControl=control.svmRadial)
> # summarize fit
> print(fit.svmRadial$results)
  sigma      C Accuracy      Kappa AccuracySD      KappaSD
1 0.04419447 0.25 0.8672520 0.2457674 0.01803261 0.1117821
2 0.04419447 0.50 0.8996927 0.4850752 0.03357527 0.1936929
3 0.04419447 1.00 0.9026778 0.5334358 0.03377470 0.1934882

```

#3. random forest

```

> control.rf = trainControl(method="cv", number=5)
> set.seed(7)
> metric = "Accuracy"
> mtry = 7 # mtry=7 (number of variables to try)
> tuneGrid <- expand.grid(.mtry=mtry)
> fit.rf_default <- train(glyhb~., data=trainingData, method="rf", metric=metric, tuneGrid=tuneGrid, preProc=c("center", "scale"), trControl=control)
> print(fit.rf_default$results)
> #4. parameter tuning via grid search for random forest
> control.rf_search <- trainControl(method="repeatedcv", number=5, repeats=3, search="grid")
> set.seed(7)
> tuneGrid <- expand.grid(.mtry=c(1:15))
> fit.rf_gridsearch <- train(glyhb~., data=trainingData, method="rf", metric=metric, tuneGrid=tuneGrid, trControl=control.rf_search, ntree=1000)
> print(fit.rf_gridsearch) # accuracy = 0.9204355 when mtry = 12

```

Random Forest


```
339 samples
14 predictor
2 classes: '0', '1'
```

No pre-processing

Resampling: Cross-validated (5 fold, repeated 3 times)

Summary of sample sizes: 271, 271, 272, 271, 271, 272, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
1	0.8535218	0.1258378
2	0.8987098	0.5300869
3	0.9144692	0.6339489
4	0.9193858	0.6581317
5	0.9222981	0.6750889
6	0.9222835	0.6785513
7	0.9242447	0.6863912
8	0.9232643	0.6838630
9	0.9232497	0.6832487
10	0.9193277	0.6721862
11	0.9163719	0.6622605
12	0.9134449	0.6507448
13	0.9114841	0.6445033
14	0.9124495	0.6516127
15	0.9065379	0.6273732

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 7.

```
> print(fit.rf_gridsearch$finalModel)
```

Call:

```
randomForest(x = x, y = y, ntree = 1000, mtry = param$mtry)
```

Type of random forest: classification

Number of trees: 1000

No. of variables tried at each split: 7

OOB estimate of error rate: 8.55%

Confusion matrix:

0 1 class.error

0 274 12 0.04195804

1 17 36 0.32075472

```
> plot(fit.rf_gridsearch)
```

```
> # make predictions on the validation set
```

```
> set.seed(7)
```


```
> predictions = predict(fit.rf_gridsearch, newdata=validationData)
```

```
> confusionMatrix = confusionMatrix(predictions, validationData$glyhb)
```

```
> # confusion matrix
```

```
> print(confusionMatrix$table)
```

	Reference	
Prediction	0	1
0	274	12
1	17	36



```
# save the final classifier model into disk
> saveRDS(fit.rf_gridsearch, "C:\\Users\\Tejal Chaudhary\\Desktop\\Exploratory-Data-Analysis-and-Prediction-on-Diabetes-Dataset-using-R-master\\diabetes_classification")
```

```
> # load the model from the disk
> final_model <- readRDS("C:\\Users\\Tejal Chaudhary\\Desktop\\Exploratory-Data-Analysis-and-Prediction-on-Diabetes-Dataset-using-R-master\\diabetes_classification")
> #final_model <- readRDS("C:\\data\\diabetes_classification")
> print(final_model)
Random Forest
```

```
339 samples
14 predictor
2 classes: '0', '1'
```

```
No pre-processing
Resampling: Cross-validated (5 fold, repeated 3 times)
Summary of sample sizes: 271, 271, 272, 271, 271, 272, ...
Resampling results across tuning parameters:
```

mtry	Accuracy	Kappa
1	0.8535218	0.1258378
2	0.8987098	0.5300869
3	0.9144692	0.6339489
4	0.9193858	0.6581317
5	0.9222981	0.6750889
6	0.9222835	0.6785513
7	0.9242447	0.6863912
8	0.9232643	0.6838630
9	0.9232497	0.6832487
10	0.9193277	0.6721862
11	0.9163719	0.6622605
12	0.9134449	0.6507448
13	0.9114841	0.6445033
14	0.9124495	0.6516127
15	0.9065379	0.6273732

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mtry = 7.

```
# make predictions using the loaded model
> set.seed(7)
> predictions = predict(final_model, newdata=validationData)
> print(predictions)
[1] 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0
0
Levels: 0 1
```